
Pandas Exam Paper 2 - (Total Marks 30 Questions - 2 Marks Each)

Section A: Data Manipulation (7 Questions)

1. Applying Functions to Columns

Apply a function to double the values of the 'Price' column using `apply()` .

Answer: `import pandas as pd`

Sample DataFrame

```
data = {  
    'Product': ['A', 'B', 'C'], 'Price': [10, 20, 30]  
}  
  
df = pd.DataFrame(data)
```

Applying function to double the 'Price' column

```
df['Price'] = df['Price'].apply(lambda x: x * 2)
```

```
print(df)
```

Output:

	Product	Price
0	A	20
1	B	40
2	C	60

2. Mapping Values in Series

Use `map()` to replace all occurrences of 'Yes' in the 'Passed' column with `True` and 'No' with `False` .

Answer: `import pandas as pd`

```
# Sample DataFrame
```

```
data = {'Student': ['Alice', 'Bob', 'Charlie'], 'Passed': ['Yes', 'No', 'Yes']}
```

```
df = pd.DataFrame(data)
```

```
# Mapping values in 'Passed' column
```

```
df['Passed'] = df['Passed'].map({'Yes': True, 'No': False})
```

```
print(df)
```

```
Output : Student Passed
```

```
0 Alice True
```

```
1 Bob False
```

```
2 Charlie True
```

3. Lowercase Strings

Convert all strings in the 'Names' column to lowercase.

Answer: import pandas as pd

```
# Sample DataFrame
```

```
data = {'Names': ['Alice', 'BOB', 'ChArLie']}
```

```
df = pd.DataFrame(data)
```

```
# Converting all names to lowercase
```

```
df['Names'] = df['Names'].str.lower()
```

```
print(df)
```

```
Output : Names
```

```
0 alice
```

```
1 bob
```

```
2 charlie
```

4. Uppercase Strings

Convert the 'City' column to uppercase.

Answer: import pandas as pd

Sample DataFrame

data = {'City': ['New York', 'los angeles', 'chicago']}

df = pd.DataFrame(data)

Converting all city names to uppercase

df['City'] = df['City'].str.upper()

print(df)

Output : City

0 NEW YORK

1 LOS ANGELES

2 CHICAGO

5. Splitting Strings

Split the 'FullName' column into 'FirstName' and 'LastName' using a space as the delimiter.

Answer:

import pandas as pd

Sample DataFrame

data = {'FullName': ['Alice Johnson', 'Bob Smith', 'Charlie Brown']}

df = pd.DataFrame(data)

Splitting 'FullName' into 'FirstName' and 'LastName'

df[['FirstName', 'LastName']] = df['FullName'].str.split(' ', expand=True)

print(df)

Output : FullName FirstName LastName

0 Alice Johnson Alice Johnson

1 Bob Smith Bob Smith

2 Charlie Brown Charlie Brown

6. String Contains

Filter rows where the 'Email' column contains '@gmail.com'.

Answer: import pandas as pd

Sample DataFrame

```
data = {'Name': ['Alice', 'Bob', 'Charlie'],  
        'Email': ['alice@gmail.com', 'bob@yahoo.com', 'charlie@gmail.com']}  
df = pd.DataFrame(data)
```

Filtering rows where 'Email' contains '@gmail.com'

```
gmail_df = df[df['Email'].str.contains('@gmail.com')]
```

```
print(gmail_df)
```

Output : Name Email

0 Alice alice@gmail.com

2 Charlie charlie@gmail.com

7. Replacing String Patterns

Use `str.replace()` to replace the domain in all emails from '@example.com' to '@newdomain.com'.

Answer: import pandas as pd

Sample DataFrame

```
data = {'Name': ['Alice', 'Bob', 'Charlie'],  
        'Email': ['alice@example.com', 'bob@example.com', 'charlie@example.com']}  
df = pd.DataFrame(data)
```

Replacing '@example.com' with '@newdomain.com'

```
df['Email'] = df['Email'].str.replace('@example.com', '@newdomain.com', regex=False)
```

```
print(df)
```

Output :

	Name	Email
--	------	-------

0	Alice	alice@newdomain.com
---	-------	---------------------

1	Bob	bob@newdomain.com
---	-----	-------------------

2	Charlie	charlie@newdomain.com
---	---------	-----------------------

Section B: Grouping and Aggregation (8 Questions)

8. Grouping Data

Group the DataFrame by the 'Department' column and calculate the mean salary for each department.

Answer: `import pandas as pd`

```
# Sample DataFrame
```

```
data = {'Employee': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],
```

```
        'Department': ['HR', 'IT', 'IT', 'HR', 'Finance'],
```

```
        'Salary': [50000, 60000, 70000, 55000, 65000]}
```

```
df = pd.DataFrame(data)
```

```
# Grouping by 'Department' and calculating mean salary
```

```
dept_salary_mean = df.groupby('Department')['Salary'].mean()
```

```
print(dept_salary_mean)
```

Output :

Department	Salary
------------	--------

Finance	65000.0
---------	---------

HR	52500.0
----	---------

IT	65000.0
----	---------

Name: Salary, dtype: float64

9. Aggregating Data

Apply multiple aggregate functions (mean, max) to the 'Sales' column using `agg()` .

Answer: import pandas as pd

Sample DataFrame

```
data = {'Employee': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],  
        'Sales': [1000, 1500, 1200, 1800, 1300]}
```

```
df = pd.DataFrame(data)
```

Applying multiple aggregate functions to the 'Sales' column

```
sales_agg = df['Sales'].agg(['mean', 'max'])
```

```
print(sales_agg)
```

Output : mean 1360.0

max 1800.0

Name: Sales, dtype: float64

10. Aggregate Multiple Functions

Use `aggregate()` to calculate both the sum and count of the 'Marks' column.

Answer:

import pandas as pd

Sample DataFrame

```
data = {'Student': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],  
        'Marks': [85, 90, 78, 92, 88]}
```

```
df = pd.DataFrame(data)
```

```
# Applying multiple aggregate functions (sum and count) to the 'Marks' column
```

```
marks_agg = df['Marks'].agg(['sum', 'count'])
```

```
print(marks_agg)
```

```
Output: sum    433
```

```
count      5
```

```
Name: Marks, dtype: int64
```


11. Filtering with `isin()`

Filter rows where the 'City' column is either 'New York' or 'Los Angeles' using `isin()` .

Answer: `import pandas as pd`

Sample DataFrame

```
data = {'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],  
        'City': ['New York', 'Los Angeles', 'Chicago', 'New York', 'Houston']}  
df = pd.DataFrame(data)
```

Filtering rows where 'City' is 'New York' or 'Los Angeles'

```
filtered_df = df[df['City'].isin(['New York', 'Los Angeles'])]
```

```
print(filtered_df)
```

Output :

	Name	City
--	------	------

0	Alice	New York
---	-------	----------

1	Bob	Los Angeles
---	-----	-------------

3	David	New York
---	-------	----------

12. Grouping and Aggregating

Group the DataFrame by 'Gender' and calculate the sum of the 'Marks' column for each group.

Answer: `import pandas as pd`

Sample DataFrame

```
data = {'Student': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],  
        'Gender': ['Female', 'Male', 'Male', 'Male', 'Female'],  
        'Marks': [85, 90, 78, 92, 88]}  
df = pd.DataFrame(data)
```

Grouping by 'Gender' and calculating the sum of 'Marks'

```
marks_sum = df.groupby('Gender')['Marks'].sum()
```

```
print(marks_sum)
```

Output : Gender

Female 173

Male 260

Name: Marks, dtype: int64

13. Multiple Aggregations on Multiple Columns

Perform multiple aggregations (min, max, mean) on the 'Age' and 'Salary' columns.

Answer: import pandas as pd

Sample DataFrame

data = {'Employee': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],

'Age': [25, 30, 35, 40, 28],

'Salary': [50000, 60000, 70000, 80000, 55000]}

df = pd.DataFrame(data)

Applying multiple aggregations to 'Age' and 'Salary' columns

agg_results = df[['Age', 'Salary']].agg(['min', 'max', 'mean'])

print(agg_results)

Output : Age Salary

min 25.00 50000.0

max 40.00 80000.0

mean 31.60 63000.0

14. Grouping and Counting

Group by 'City' and count the number of entries in each city.

Answer: import pandas as pd

Sample DataFrame

data = {'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve', 'Frank'],

'City': ['New York', 'Los Angeles', 'Chicago', 'New York', 'Los Angeles', 'Chicago']}

```
df = pd.DataFrame(data)
```

```
# Grouping by 'City' and counting the number of entries
```

```
city_counts = df.groupby('City')['Name'].count()
```

```
print(city_counts)
```

Output : City

Chicago 2

Los Angeles 2

New York 2

Name: Name, dtype: int64

15. Using `apply()` with `Groupby`

Apply a custom function to find the range (max-min) of the 'Salary' column for each department.

Answer: import pandas as pd

```
# Sample DataFrame
```

```
data = {'Employee': ['Alice', 'Bob', 'Charlie', 'David', 'Eve', 'Frank'],
```

```
        'Department': ['HR', 'IT', 'IT', 'HR', 'Finance', 'Finance'],
```

```
        'Salary': [50000, 60000, 70000, 55000, 65000, 75000]}
```

```
df = pd.DataFrame(data)
```

```
# Defining a custom function to calculate the range
```

```
def salary_range(x):
```

```
    return x.max() - x.min()
```

```
# Applying the function using groupby and apply
```

```
salary_range_by_dept = df.groupby('Department')['Salary'].apply(salary_range)
```

```
print(salary_range_by_dept)
```

Output : Department

Finance 10000

HR 5000

IT 10000

Name: Salary, dtype: int64

Section C: Merging, Joining, and Concatenating (5 Questions)

16. Concatenating DataFrames

Concatenate two DataFrames `df1` and `df2` along rows.

Answer: import pandas as pd

Creating first DataFrame

```
df1 = pd.DataFrame({'ID': [1, 2, 3],  
                    'Name': ['Alice', 'Bob', 'Charlie']})
```

Creating second DataFrame

```
df2 = pd.DataFrame({'ID': [4, 5],  
                    'Name': ['David', 'Eve']})
```

Concatenating along rows

```
df_combined = pd.concat([df1, df2], axis=0, ignore_index=True)
```

```
print(df_combined)
```

Output : ID Name

0 1 Alice

1 2 Bob

2 3 Charlie

3 4 David

4 5 Eve

17. Merging DataFrames

Merge two DataFrames df1 and df2 on the 'ID' column.

Answer: import pandas as pd

Creating first DataFrame

```
df1 = pd.DataFrame({'ID': [1, 2, 3],  
                    'Name': ['Alice', 'Bob', 'Charlie']})
```

Creating second DataFrame

```
df2 = pd.DataFrame({'ID': [1, 2, 3],  
                    'Salary': [50000, 60000, 70000]})
```

Merging on 'ID' column

```
df_merged = pd.merge(df1, df2, on='ID')
```

```
print(df_merged)
```

Output : ID Name Salary

```
0 1 Alice 50000
1 2 Bob 60000
2 3 Charlie 70000
```

18. Merging with Different Keys

Merge DataFrames on different column names: 'df1' has 'EmployeeID' and 'df2' has 'ID'.

Answer: `import pandas as pd`

Creating first DataFrame

```
df1 = pd.DataFrame({'EmployeeID': [1, 2, 3],
                    'Name': ['Alice', 'Bob', 'Charlie']})
```

Creating second DataFrame

```
df2 = pd.DataFrame({'ID': [1, 2, 3],
                    'Salary': [50000, 60000, 70000]})
```

Merging on different column names

```
df_merged = pd.merge(df1, df2, left_on='EmployeeID', right_on='ID')
```

```
print(df_merged)
```

Output : `EmployeeID Name ID Salary`

```
0      1  Alice  1  50000
1      2   Bob  2  60000
2      3 Charlie  3  70000
```

19. Concatenating Along Columns

Concatenate two DataFrames df1 and df2 along columns.

Answer: `import pandas as pd`

Creating first DataFrame

```
df1 = pd.DataFrame({'ID': [1, 2, 3],  
                    'Name': ['Alice', 'Bob', 'Charlie']})
```

Creating second DataFrame

```
df2 = pd.DataFrame({'Age': [25, 30, 35],  
                    'Salary': [50000, 60000, 70000]})
```

Concatenating along columns

```
df_combined = pd.concat([df1, df2], axis=1)
```

```
print(df_combined)
```

Output :

	ID	Name	Age	Salary
0	1	Alice	25	50000
1	2	Bob	30	60000
2	3	Charlie	35	70000

20. Joining DataFrames

Join df1 and df2 on the 'ID' column with an outer join.

Answer: import pandas as pd

Creating first DataFrame

```
df1 = pd.DataFrame({'ID': [1, 2, 3],  
                    'Name': ['Alice', 'Bob', 'Charlie']})
```

Creating second DataFrame

```
df2 = pd.DataFrame({'ID': [2, 3, 4],  
                    'Salary': [60000, 70000, 80000]})
```

Performing an outer join on 'ID'

```
df_joined = pd.merge(df1, df2, on='ID', how='outer')
```

```
print(df_joined)
```

Output :

	ID	Name	Salary
--	----	------	--------

0	1	Alice	NaN
---	---	-------	-----

1	2	Bob	60000.0
---	---	-----	---------

2	3	Charlie	70000.0
---	---	---------	---------

3	4	NaN	80000.0
---	---	-----	---------

Section D: Reshaping and Input/Output (10 Questions)

21. Transposing DataFrames

Transpose the rows and columns of the DataFrame df .

Answer:

```
import pandas as pd
```

```
# Sample DataFrame
```

```
df = pd.DataFrame({'ID': [1, 2, 3],  
                   'Name': ['Alice', 'Bob', 'Charlie'],  
                   'Salary': [50000, 60000, 70000]})
```

```
# Transposing the DataFrame
```

```
df_transposed = df.T
```

```
print(df_transposed)
```

Output :

	1	2
--	---	---

ID	1	2	3
----	---	---	---

Name	Alice	Bob	Charlie
------	-------	-----	---------

Salary 50000 60000 70000

22. Using T Attribute

Use the T attribute to transpose the DataFrame df .

Answer: import pandas as pd

Sample DataFrame

```
df = pd.DataFrame({'ID': [1, 2, 3],  
                  'Name': ['Alice', 'Bob', 'Charlie'],  
                  'Salary': [50000, 60000, 70000]})
```

Transposing the DataFrame using .T

```
df_transposed = df.T
```

```
print(df_transposed)
```

```
Output :      0    1    2
```

```
ID      1    2    3
```

```
Name Alice  Bob  Charlie
```

```
Salary 50000 60000 70000
```

23. Writing to CSV

Save the DataFrame df to a file called output.csv .

Answer: import pandas as pd

Sample DataFrame

```
df = pd.DataFrame({'ID': [1, 2, 3],  
                  'Name': ['Alice', 'Bob', 'Charlie'],  
                  'Salary': [50000, 60000, 70000]})
```

Saving DataFrame to CSV

```
df.to_csv('output.csv', index=False)
```

```
print("DataFrame saved to output.csv")
```

24. Writing to Excel

Export the DataFrame `df` to an Excel file named `output.xlsx`.

Answer: `import pandas as pd`

Sample DataFrame

```
df = pd.DataFrame({'ID': [1, 2, 3],  
                  'Name': ['Alice', 'Bob', 'Charlie'],  
                  'Salary': [50000, 60000, 70000]})
```

Saving DataFrame to Excel

```
df.to_excel('output.xlsx', index=False, engine='openpyxl')
```

```
print("DataFrame saved to output.xlsx")
```

25. Writing to JSON

Convert the DataFrame `df` to a JSON file named `output.json`.

Answer: `import pandas as pd`

Sample DataFrame

```
df = pd.DataFrame({'ID': [1, 2, 3],  
                  'Name': ['Alice', 'Bob', 'Charlie'],  
                  'Salary': [50000, 60000, 70000]})
```

Saving DataFrame to JSON

```
df.to_json('output.json', orient='records', indent=4)
```

```
print("DataFrame saved to output.json")
```

26. Rendering DataFrame as HTML

Convert the DataFrame `df` to an HTML table and save it as `output.html`.

Answer: `import pandas as pd`

```
# Sample DataFrame
```

```
df = pd.DataFrame({'ID': [1, 2, 3],  
                   'Name': ['Alice', 'Bob', 'Charlie'],  
                   'Salary': [50000, 60000, 70000]})
```

27. Loading CSV File

Load a CSV file named `student_data.csv` into a DataFrame.

Answer:

```
import pandas as pd
```

```
# Loading the CSV file into a DataFrame
```

```
df = pd.read_csv('student_data.csv')
```

```
# Display the first few rows
```

```
print(df.head())
```

```
df = pd.read_csv('/path/to/student_data.csv')
```

29. Saving a DataFrame as

28. Loading Excel File

Load an Excel file named `sales_data.xlsx` into a DataFrame.

Answer: `import pandas as pd`

Loading the Excel file into a DataFrame

```
df = pd.read_excel('sales_data.xlsx', engine='openpyxl')
```

Display the first few rows

```
print(df.head())
```

```
df = pd.read_excel('sales_data.xlsx', sheet_name='Sheet1', engine='openpyxl')
```

```
df = pd.read_excel('/path/to/sales_data.xlsx', engine='openpyxl')
```

29. Saving a DataFrame as CSV

Save the DataFrame `df` to a CSV file called `employees.csv`, including only the 'Name' and 'Salary' columns.

Answer: `import pandas as pd`

Sample DataFrame

```
df = pd.DataFrame({'ID': [1, 2, 3],  
                  'Name': ['Alice', 'Bob', 'Charlie'],  
                  'Salary': [50000, 60000, 70000],  
                  'Department': ['HR', 'IT', 'Finance']})
```

Saving only 'Name' and 'Salary' columns to CSV

```
df[['Name', 'Salary']].to_csv('employees.csv', index=False)
```

```
print("DataFrame saved to employees.csv")
```

Output : Name,Salary

Alice,50000

Bob,60000

Charlie,70000

30. Saving a DataFrame as JSON with Specific Columns

Save the DataFrame `df` as a JSON file, but only include the 'Name' and 'Department' columns.

Answer: `import pandas as pd`

Sample DataFrame

```
df = pd.DataFrame({'ID': [1, 2, 3],  
                  'Name': ['Alice', 'Bob', 'Charlie'],  
                  'Salary': [50000, 60000, 70000],  
                  'Department': ['HR', 'IT', 'Finance']})
```

Saving only 'Name' and 'Department' columns to JSON

```
df[['Name', 'Department']].to_json('employees.json', orient='records', indent=4)
```

```
print("DataFrame saved to employees.json")
```
