

# Reinforcement Learning

MUKESH KUMAR

# AGENDA

- What is Reinforcement Learning
- Core Concepts of RL
- Types of RL
- Q-Learning Algorithm
- Mouse Maze Problem
- Q-Learning in OpenAI Gym – Smart Taxi
- SARSA
- Policy Gradient Theory

# Introduction to Reinforcement Learning

Reinforcement Learning (RL) is a type of machine learning where an agent learns to make decisions by interacting with an environment to maximize a reward.

# Introduction to Reinforcement Learning

**Definition:** Learning by interacting with an environment to maximize cumulative reward.

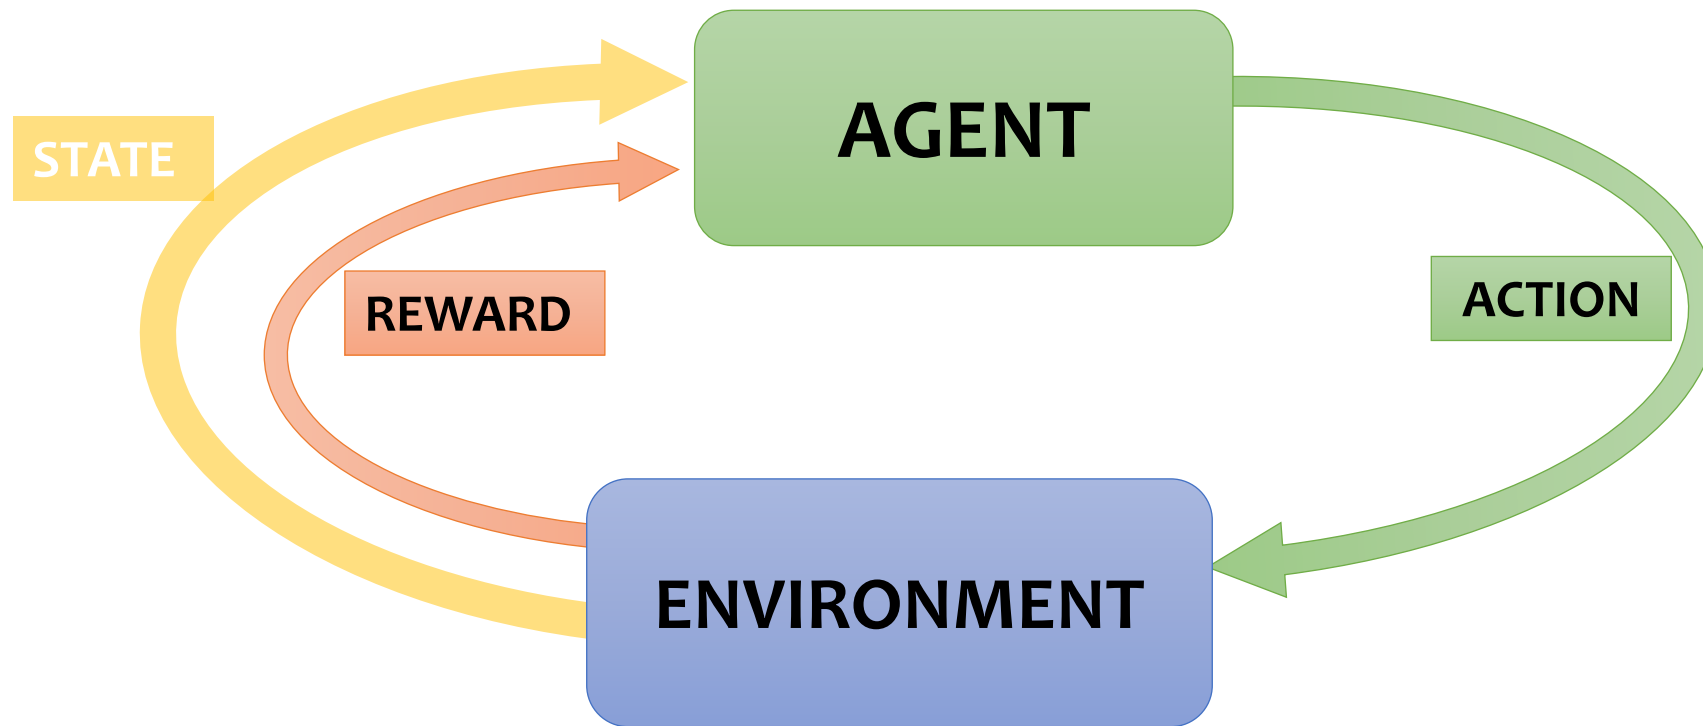
## Key Concepts:

- **Agent:** The learner or decision-maker.
- **Environment:** The world the agent interacts with.
- **Action:** A move the agent makes.
- **State:** The current situation of the agent.
- **Reward:** Feedback signal from the environment.

## Goal:

Learn a **policy** (strategy) that maximizes the **cumulative reward** over time.

# Architecture of Reinforcement Learning



# Where is Reinforcement Learning Used?

## **Games & Simulation:**

- AlphaGo (beating world champions at Go)
- OpenAI Five (Dota 2)
- Chess and Atari games

## **Robotics:**

- Robotic arms for picking/sorting
- Industrial automation
- Bipedal/humanoid robot walking

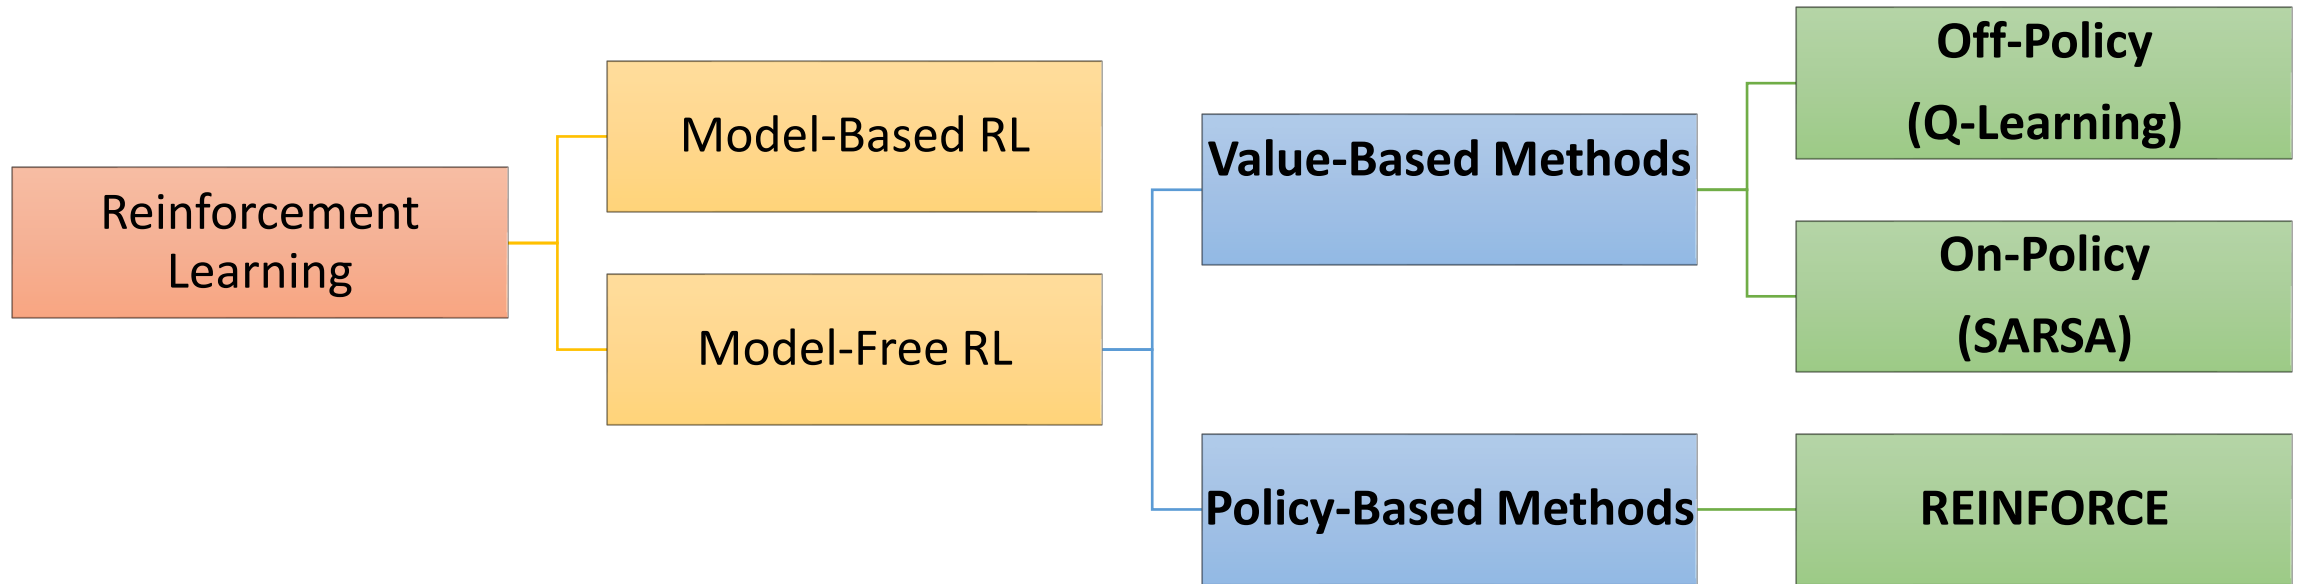
# Where is Reinforcement Learning Used?

## **Autonomous Driving:**

- Lane following
- Obstacle avoidance
- Decision-making in traffic

## **Operations & Control:**

- Elevator scheduling
- HVAC systems
- Power grid optimization





# Key Components of Reinforcement Learning

## Agent

- The decision maker
- Learns from experience and chooses actions

## Environment

- Everything the agent interacts with
- Provides feedback via rewards

# Key Components of Reinforcement Learning

## State (s)

- A snapshot of the environment at a given time
- Example: Position of a car in a driving simulation

## Action (a)

- A decision taken by the agent
- Example: Move left, accelerate, pick object


# Key Components of Reinforcement Learning

## Reward ( $r$ )

- A numerical signal from the environment
- Positive for good actions, negative for bad ones
- Drives learning and behavior

# Mouse Maze

## Problem Introduction




- There are totally 10 States = 0-9
- For each state what are the possible actions and reward for each action( reward can be positive or negative)
- Goal: Mouse has to exit the maze

# Setting up the Rewards

- Reward is a scalar quantity
  - Positive
  - Negative
  - Neutral

For mouse maze , we only have positive rewards, there is not Negative rewards



# Initializing Matrices

- Q-Matrix
  - R-Matrix
  - State-Action Reward Matrix
- 
- What are the different starts in mouse maze problem?
  - What are the actions possible in mouse maze problems??



# Reward Matrix

- 10 States possible = 0-9
- 4 actions possible = up, down, right, left

	up	down	left	right
0				
1				
2				
3				
4				
5				
6				
7				
8				
9				

# For state 0

- From 0 it can only go to 3 , only down action is possible
- From State 1, only right is possible

	up	down	left	right
0		0		
1				0
2				
3				
4				
5				
6				
7				
8				
9				

# Another approach of creating Reward Matrix

- State vs State

We are going to use this one for  
Mouse maze problem

[illegible]

# Sate Vs State Reward Matrix

- From 0 state it can move to only 3
- From 1 state it can move only to 2
- From 2 state it can go to 1,5 or 9
- This is what we call reward matrix, what are the possible actions and corresponding rewards

[illegible]

# Reward Matrix

- This is the final Reward matrix , -1 here is not reward it's an invalid option for python program

-1	-1	-1	0	-1	-1	-1	-1	-1	-1	-1
-1	-1	0	-1	-1	-1	-1	-1	-1	-1	-1
-1	0	-1	-1	-1	-1	0	-1	-1	-1	100
0	-1	-1	-1	0	-1	0	-1	-1	-1	-1
-1	-1	-1	0	-1	0	-1	0	-1	-1	-1
-1	-1	0	-1	0	-1	-1	-1	-1	-1	-1
-1	-1	-1	0	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	0	-1	-1	-1	-1	0	100
-1	-1	-1	-1	-1	-1	-1	-1	0	-1	-1
-1	-1	0	-1	-1	-1	-1	-1	0	-1	100



# Step1: Calculation

- Lets say we drop the mouse at 2
- From 2 : 1, 5, 9
- Say it goes to 9
- $Q(2,9) = R(2,9) + 0.7 ( Q(9,2), Q(9,9), Q(9,7)) - Q(2,9)$
- $Q(2,9) = 100$

# Updated Q-Matrix after step 1



- When we repeat this process 1000 of times we will have a Q-matrix which will have optimum values in each cell
- We can then use this Q-Matrix to decide the next best state mouse should move in order to exit the maze in shortest possible path
- $0 > 3 > 4 > 7 > 9$

[illegible]



- In Class Problem

Calculate the Q value, when the initial state = 3 and mouse moves to state 4 at Random. (Assume Gamma = 0.7)



# Q-Learning Algorithm (Theory)

- **Q-Learning** is a model-free, value-based **Reinforcement Learning algorithm**
- that learns the **optimal action-selection policy** using **Q-values**.
  
- Learns how good an action is at a given state, without knowing the environment's
- dynamics.

# Q-Learning Key Concepts

## Q-Value ( $Q(s, a)$ ):

- Estimates the **expected total reward** of taking action **a** in state **s** and following the optimal policy afterward.

## Goal:

- Learn the optimal Q-values, denoted  $Q^*$ , so the agent can act greedily:

$$\pi(s) = \operatorname{argmax}_a Q^*(s, a)$$

# Q-Learning Update Rule


$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[ r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

- Where:
- **s** = current state
- **a** = action taken
- **r** = reward received
- **s'** = next state
- **a'** = possible next actions
- **α** = learning rate ( $0 < \alpha \leq 1$ )
- **γ** = discount factor ( $0 \leq \gamma < 1$ )

$$Q(s, a) \leftarrow Q(s, a) + \alpha \cdot \underbrace{\left[ r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]}_{\text{TD Error}=\delta}$$

Temporal Difference (TD)

## Q-learning: Algorithm



# Characteristics of Q-Learning

## **Model-Free**

- Doesn't require knowledge of transition probabilities

## **Off-Policy**

- Learns optimal policy independently of the actions taken

## **Convergence**

- Converges to optimal Q-values under appropriate exploration and learning rate conditions